

Vector-attributed graph G and threshold τ



Phase 1: pivoted structural BK
 τ -independent enumeration of each structural maximal clique M



Top-Edge bound
certifies size layers

Vertex-Incidence bound
tightens layer feasibility



Phase 2: exact descending subset enumeration
test only certified layers; skip subsets contained in a feasible local superset

Certified pruning
preserves exact output



Candidate pool $\mathcal{F}$ from all structural cliques



Phase 3: exact duplicate removal and global strict-containment filtering



Inclusion-wise maximal ACMSCs $\mathcal{R}$